

Data Immersion
Databases & SQL for Analysts
3.6: Summarizing & Cleaning Data in SQL
David Guillen Aroche

1. Check for and clean dirty data for the film table and the customer table.

- Duplicates.

Film Table

Query:

```
SELECT film_id,  
       title,  
       description,  
       release_year,  
       language_id,  
       rental_duration,  
       rental_rate,  
       film.length,  
       replacement_cost,  
       rating,  
       last_update,  
       special_features,  
       fulltext,  
       COUNT(*)  
FROM film  
GROUP BY film_id,  
         title,  
         description,  
         release_year,  
         language_id,  
         rental_duration,  
         rental_rate,  
         film.length,  
         replacement_cost,  
         rating,  
         last_update,  
         special_features,  
         fulltext  
HAVING COUNT(*) >1
```

Result:

No duplicates were found in film table, see screenshot below.



ase_year	language_id	rental_duration	rental_rate	length	replacement_cost	rating	last_update	special_features	fulltext	count
ger	smallint	smallint	numeric (4,2)	smallint	numeric (5,2)	mpaa_rating	timestamp without time zone	text[]	tsvector	bigint

Total rows: 0 of 0 Query complete 00:00:00.202 Ln 254, Col 19

In case of finding duplicate values, I would have created a view with unique records with the following syntax to make each row unique:

```
CREATE VIEW viewname AS
SELECT col1,
       col2,
       col3 ...
FROM tablename
GROUP BY col1,
         col2,
         col3, ... ;
```

Customer Table

Query:

```
SELECT customer_id,
       store_id,
       first_name,
       email,
       address_id,
       activebool,
       last_update,
       active,
       COUNT(*)
FROM customer
GROUP BY customer_id,
         store_id,
         first_name,
         email,
         address_id,
         activebool,
```

```

last_update,
active
HAVING COUNT(*) > 1

```

Result:

No duplicates were found in customers table, see screenshot below.

The screenshot shows a database query interface. At the top, there's a toolbar with icons for query execution, saving, and other functions. Below the toolbar, a table schema is displayed with columns: `ase_year` (type: ger), `language_id` (type: smallint), `rental_duration` (type: smallint), `rental_rate` (type: numeric (4,2)), `length` (type: smallint), `replacement_cost` (type: numeric (5,2)), `rating` (type: mpaa_rating), `last_update` (type: timestamp without time zone), `special_features` (type: text[]), `fulltext` (type: tsvector), and `count` (type: bigint). Below the schema, a status bar indicates "Total rows: 0 of 0" and "Query complete 00:00:00.202". On the right side of the status bar, it says "Ln 254, Col 19".

In case of finding duplicate values, I would have created a view with unique records with the following syntax to make each row unique:

```

CREATE VIEW viewname AS
SELECT col1,
       col2,
       col3 ...
FROM tablename
GROUP BY col1,
         col2,
         col3, ... ;

```

- **Non-uniform data.**

Film Table

Query (syntax):

```

SELECT DISTINCT column_name
FROM ctable_name
GROUP BY column_name

```

Results:

Non-uniform values were searched in the table, but no irregularities were found. The total number of values for each column is as follows:

- *film_id, 1000 values*
- *title, 1000 values*
- *description, 1000 values*
- *release_year, 1 value*

- *language_id*, 1 value
- *rental_duration*, 5 values
- *rental_rate*, 3 values
- *film.length*, 140 values
- *replacement_cost*, 21 values
- *rating*, 5 values
- *last_update*, 1 value
- *special_features*, 15 vlaues
- *fulltext*, 1000 values

In case of finding a non-uniform value, first, I would have grouped the values to get familiar with them to choose a standard format, and checked on data types for each column. Then use the following syntax to update values:

```
UPDATE film
SET rating = 'G'
WHERE rating IN ('gen',
                'g',
                'General')
```

Customer Table

Query (syntax):

```
SELECT DISTINCT column_name
FROM ctable_name
GROUP BY column_name
```

Results:

No Non-uniform values were searched and have found no irregularities. The total number of values for each column as follows:

- *customer_id*, 599 values
- *store_id*, 2 values
- *first_name*, 591 values
- *last_name*, 599 values
- *email*, 599 values
- *address_id*, 599 values
- *activebool*, 1 value
- *create_date*, 1 value
- *active*, 1 value

e.g.:

304	
305	SELECT DISTINCT active
306	FROM customer
307	GROUP BY active

Data output	Messages	Notification
-------------	----------	--------------

--	--	--	--	--	--	--

	active integer	
1		0
2		1

Save results to file
 F8

In case of finding a non-uniform value, first, I would have grouped the values to get familiar with them to choose a standard format, and checked on data types for each column. Then use the following syntax to update values:

```
UPDATE film
SET rating = 'G'
WHERE rating IN ('gen',
                'g',
                'General')
```

- Missing values

Film Table

While looking for **Non-uniform** values, no missing values were found.

In case of have found missing values, I would have checked on the proportion of the missing data for each column to determine if would have been valid to include them in the analysis or not, and adapt the analysis to mitigate the impact as much as possible.

Customer Table

While looking for **Non-uniform** values, no missing values were found.

In case of have found missing values, I would have checked on the proportion of the missing data for each column to determine if would have been valid to include them in the analysis or not, and adapt the analysis to mitigate the impact as much as possible.

2. Summarize Data.

Film Table

Query (Syntax) Numeric Columns:

```
SELECT MIN(column_name) AS column_alias,  
       MAX(column_name) AS column_alias,  
       AVG(column_name) AS column_alias,  
       COUNT(column_name) AS column_alias,  
       COUNT(*) AS column_alias  
FROM table_name
```

--film_id.

The average was not included for this query as it does not really say much.

```
296 SELECT MIN(film_id) AS lowest_id,  
297        MAX(film_id) AS highest_id,  
298        COUNT(film_id) AS films_count,  
299        COUNT(*) AS tuples_count  
300 FROM film
```

Data output Messages Notifications					
	lowest_id integer	highest_id integer	films_count bigint	tuples_count bigint	
1	1	1000	1000	1000	

--release_year

The average **years count** was not included since there is only one value for this column.

```
296 SELECT MIN(release_year) AS minimum_year,  
297        MAX(release_year) AS maximum_year,  
298        COUNT(release_year) AS years_count,  
299        COUNT(*) AS tuples_count  
300 FROM film
```

Data output Messages Notifications					
	minimum_year integer	maximum_year integer	years_count bigint	tuples_count bigint	
1	2006	2006	1000	1000	

--language_id

There is only one language value for the total tuples in the table.

```
296 SELECT language_id,  
297         MIN(language_id) AS minimum_language_id,  
298         MAX(language_id) AS maximum_language_id,  
299         AVG (language_id) AS average_languages_count,  
300         COUNT(language_id) AS languages_count,  
301         COUNT(*) AS tuples_count  
302 FROM film  
303 GROUP BY language_id
```

Data output Messages Notifications

	language_id smallint	minimum_language_id smallint	maximum_language_id smallint	average_languages_count numeric	languages_count bigint	tuples_count bigint
1	1	1	1	1	1000	1000

--rental_duration

```
296 SELECT MIN(rental_duration) AS minimum_rental_duration,  
297         MAX(rental_duration) AS maximum_rental_duration,  
298         AVG (rental_duration) AS average_rental_duration,  
299         COUNT(rental_duration) AS rental_duration_values_count,  
300         COUNT(*) AS tuples_count  
301 FROM film
```

Data output Messages Notifications

	minimum_rental_duration smallint	maximum_rental_duration smallint	average_rental_duration numeric	rental_duration_values_count bigint	tuples_count bigint
1	3	7	4.985	1000	1000

--rental_rate

```
296 SELECT MIN(rental_rate) AS minimum_rental_rate,  
297         MAX(rental_rate) AS maximum_rental_rate,  
298         AVG (rental_rate) AS average_rental_rate,  
299         COUNT(rental_rate) AS rental_rate_values_count,  
300         COUNT(*) AS tuples_count  
301 FROM film
```

Data output Messages Notifications

	minimum_rental_rate numeric	maximum_rental_rate numeric	average_rental_rate numeric	rental_rate_values_count bigint	tuples_count bigint
1	0.99	4.99	2.98	1000	1000

--film.length

```
296 SELECT MIN(film.length) AS minimum_film_length,  
297         MAX(film.length) AS maximum_film_length,  
298         AVG (film.length) AS average_film_length,  
299         COUNT(film.length) AS film_length_count,  
300         COUNT(*) AS tuples_count  
301 FROM film
```

Data output Messages Notifications

	minimum_film_length smallint	maximum_film_length smallint	average_film_length numeric	film_length_count bigint	tuples_count bigint
1	46	185	115.272	1000	1000

--replacement_cost

```
296 SELECT MIN(replacement_cost) AS minimum_replacement_cost,  
297         MAX(replacement_cost) AS maximum_replacement_cost,  
298         AVG (replacement_cost) AS average_replacement_cost,  
299         COUNT(replacement_cost) AS replacement_cost_count,  
300         COUNT(*) AS tuples_count  
301 FROM film
```

Data output Messages Notifications

	minimum_replacement_cost numeric	maximum_replacement_cost numeric	average_replacement_cost numeric	replacement_cost_count bigint	tuples_count bigint
1	9.99	29.99	19.984	1000	1000

Query (Syntax) Non-Numeric Columns:

```
SELECT mode() WITHIN GROUP (ORDER BY column_name)  
        AS modal_value  
FROM tablename
```

--title

```
305 SELECT mode() WITHIN GROUP (ORDER BY title)  
306         AS modal_value  
307 FROM film
```

Data output Messages Notifications

	modal_value character varying
1	Academy Dinosaur

--description

```
305 SELECT mode() WITHIN GROUP (ORDER BY description)
306     AS modal_value
307 FROM film
```

Data output Messages Notifications

	modal_value text	
1	A Action-Packed Character Study of a Astronaut And a Explorer who must Reach a Monkey in A MySQL Convention	

--rating

```
305 SELECT mode() WITHIN GROUP (ORDER BY rating)
306     AS modal_value
307 FROM film
```

Data output Messages Notifications

	modal_value mpaa_rating	
1	PG-13	

--special_features

```
305 SELECT mode() WITHIN GROUP (ORDER BY special_features)
306     AS modal_value
307 FROM film
```

Data output Messages Notifications

	modal_value text[]	
1	{Trailers,Commentaries,"Behind the Scenes"}	

--fulltext

```
305 SELECT mode() WITHIN GROUP (ORDER BY fulltext)
306     AS modal_value
307 FROM film
```

Data output Messages Notifications

	modal_value tsvector	
1	'balloon':19 'confront':14 'documentari':5 'feminist':8,11,16 'mile':2 'must':13 'spi':1 'thrill':4	

Customer Table

Query (Syntax) Numeric Columns:

```
SELECT MIN(column_name) AS column_alias,  
       MAX(column_name) AS column_alias,  
       AVG(column_name) AS column_alias,  
       COUNT(column_name) AS column_alias,  
       COUNT(*) AS column_alias  
FROM table_name
```

--customer_id

```
312 SELECT MIN(customer_id) AS minimum_customer_id,  
313        MAX(customer_id) AS maximum_customer_id,  
314        AVG (customer_id) AS average_customer_ids,  
315        COUNT(customer_id) AS customer_id_count,  
316        COUNT(*) AS tuples_count  
317 FROM customer
```

Data output Messages Notifications

	minimum_customer_id integer	maximum_customer_id integer	average_customer_ids numeric	customer_id_count bigint	tuples_count bigint
1	1	599	300	599	599

--storage_id

```
312 SELECT MIN(store_id) AS minimum_store_id,  
313        MAX(store_id) AS maximum_store_id,  
314        AVG (store_id) AS average_store_ids,  
315        COUNT(store_id) AS store_id_count,  
316        COUNT(*) AS tuples_count  
317 FROM customer
```

Data output Messages Notifications

	minimum_store_id smallint	maximum_store_id smallint	average_store_ids numeric	store_id_count bigint	tuples_count bigint
1	1	2	1.455759599332220	599	599

--address_id

```
312 SELECT MIN(address_id) AS minimum_address_id,  
313         MAX(address_id) AS maximum_address_id,  
314         AVG (address_id) AS average_address_ids,  
315         COUNT(address_id) AS address_id_count,  
316         COUNT(*) AS tuples_count  
317 FROM customer
```

Data output Messages Notifications

	minimum_address_id smallint	maximum_address_id smallint	average_address_ids numeric	address_id_count bigint	tuples_count bigint
1	5	605	304.7245409015025	599	599

--create_date

Not possible to get average for dates.

```
312 SELECT MIN(create_date) AS minimum_create_date,  
313         MAX(create_date) AS maximum_create_date,  
314         COUNT(create_date) AS create_date_count,  
315         COUNT(*) AS tuples_count  
316 FROM customer
```

Data output Messages Notifications

	minimum_create_date date	maximum_create_date date	create_date_count bigint	tuples_count bigint
1	2006-02-14	2006-02-14	599	599

--last_update

Not possible to get averages for dates.

```
312 SELECT MIN(last_update) AS minimum_last_update,  
313         MAX(last_update) AS maximum_last_update,  
314         COUNT(last_update) AS create_last_update,  
315         COUNT(*) AS tuples_count  
316 FROM customer
```

Data output Messages Notifications

	minimum_last_update timestamp without time zone	maximum_last_update timestamp without time zone	create_last_update bigint	tuples_count bigint
1	2013-05-26 14:49:45.738	2013-05-26 14:49:45.738	599	599

--active

Averages was not calculated for this columns as does not make sense to get it; only to values.

```
312 SELECT MIN(active) AS minimum_active,  
313         MAX(active) AS maximum_active,  
314         COUNT(active) AS create_active,  
315         COUNT(*) AS tuples_count  
316 FROM customer
```

Data output Messages Notifications

	minimum_active integer	maximum_active integer	create_active bigint	tuples_count bigint
1	0	1	599	599

Query (Syntax) Non-Numeric Columns:

```
SELECT mode() WITHIN GROUP (ORDER BY column_name)
```

```
AS modal_value
```

```
FROM tablename
```

--first_name

```
320 SELECT mode() WITHIN GROUP (ORDER BY first_name)  
321         AS modal_value  
322 FROM customer
```

Data output Messages Notifications

	modal_value character varying
1	Jamie

--last_name

```
320 SELECT mode() WITHIN GROUP (ORDER BY last_name)  
321         AS modal_value  
322 FROM customer
```

Data output Messages Notifications

	modal_value character varying
1	Abney

--email

Please consider there are 599, and 599 different emails. This results is for one of those 599 different emails.

```
320 SELECT mode() WITHIN GROUP (ORDER BY email)
321     AS modal_value
322 FROM customer
```

Data output Messages Notifications



	modal_value character varying
1	aaron.selby@sakilacustomer.org

3. Which tool (Excel or SQL) is more efficient for data profiling and why? Consider functions, ease of use, and speed.

In general, SQL is way more effective at providing data types, data integrity, and calculating descriptive statistics of the data. Personally, I find the difference essentially in the capacity that both have in terms of processing data, and how they show the data; SQL's results are succinct, and is considerably quicker coding a query rather than making the data profile from scratch in Excel—though making a data profile in Excel is slower and has less data processing potency, it also is more intuitive once you know what you are doing, and there are more resources around to troubleshoot in excel than in SQL.

Functions.

I believe SQL is more effective at providing results; the used functions are basically the same as in Excel, but the one difference would be knowing how to code queries with everything at once in SQL, rather than making a longer process in Excel for each one of the functions—leaving aside the data processing capacity. The downside of SQL is that you must know how to query, you have to have the technical knowledge to do so, otherwise it might be difficult to impossible to retrieve the data the way you want it.

Ease of Use.

It's easier to use Excel even though the data integrity analysis represents more time. The issue with Excel is that it has less capacity to process data. On the other hand, once you have (at least the basic) technical knowledge to code queries, it's easier to use SQL, again, because of how succinct the process and the results are

Speed.

SQL is by far the fastest way to make a data integrity analysis. It's faster to code a query, even if it's a long and/or complex one (once you have the knowledge).

In conclusion, SQL is better to make a data integrity analysis, but not the easiest.